



QKV fused    BF16 merged    INT4 + BF16 csamv LoRA    O projection



UDA Graph median latency (Decode batch / token rows) is prewarm, and L2 flush. Kernel-only Decode batch / token rows are not modeled.