


R1 7B: INT4 Base-Only Reference vs BF16 Merged



UDA Graph median latency for decode batch size of 1, 4, 8, 16, 32, 64, 128, 256, and 512. Kernel-only latencies are not modeled.